

Affidavit

Ich, Jan Monschke, geboren am 12.03.1987, versichere, diese Bachelorarbeit selbstständig und lediglich unter Benutzung der angegebenen Quellen und Hilfsmittel verfasst zu haben.

Ich erkläre weiterhin, dass die vorliegende Arbeit noch nicht im Rahmen eines anderen Prüfungsverfahrens eingereicht wurde.

Düsseldorf, den

Jan Monschke

Table of Contents

AFFIDAVIT	1
TABLE OF CONTENTS	2
I SALON	3
1. OVERVIEW	3
2. THE IDEA	3
3. IMPLEMENTATION	3
3.1 <i>Technology</i>	3
3.2 <i>Data Model</i>	4
3.3 <i>Navigation</i>	5
3.4 <i>Menu</i>	5
3.5 <i>Pages</i>	6
3.5.1 Index	6
3.5.2 Registration and Account	6
3.5.3 User Overview	7
3.5.4 Page Overview	8
3.5.5 Page Edit Form	8
3.5.6 Image Overview	9
3.5.7 Image Edit Form	9
3.5.8 Page Index	10
3.5.9 Search	11
3.6 <i>Drag and Drop</i>	11
3.7 <i>Visual Design</i>	13
4 QUO VADIS SALON?	14
5 EVALUATION	16
II SINGLE PAGE WEB APPS	16
1. INTRODUCTION	16
1.1 <i>Motivation for Salon</i>	16
1.2 <i>Comparison</i>	17
1.2.1 <i>Traditional</i>	18
1.2.2 <i>Single Page</i>	18
2 DIFFERENCES	19
2.1 <i>URLs</i>	19
2.2 <i>Rendering</i>	20
2.3 <i>Authentication</i>	20
2.4 <i>Internationalization</i>	21
2.5 <i>Notifications</i>	21
2.6 <i>Forms</i>	22
3 BENEFITS	23
3.1 <i>Speed / Efficiency</i>	23
3.2 <i>Same Language</i>	23
3.3 <i>User Experience</i>	24
4 PROBLEMS	26
4.1 <i>Search Engine Optimization</i>	26
4.2 <i>New Tools needed</i>	27
4.3 <i>Exposure of Business Logic / Sensitive Data</i>	28
4.4 <i>Accessibility</i>	28
5. CONCLUSION	29

I Salon

1. Overview

Salon is a web-based system that allows its users to create pages and to upload images onto these pages to create an online portfolio of their work. On a first sight this functionality may not look very innovative since there are many other services on the Internet that allow the user to upload images. But the main improvement that Salon offers that other services don't offer is that registered users are able to fully control the way their images are presented to their visitors.

Other services, like for example Flickr, don't let users decide about the presentation of their work and only provide simple list views and slide shows. In that way all albums look the same.

In Salon all images are placed on a canvas and can freely get dragged around by the user to create innovative and unique arrangements. Also the canvas itself can be moved to focus a certain point of a page. Another feature is that images can link to other pages so that users can create associations between pages or even associations between users. These features give the user another way to express creativity.

2. The idea

Dipl. Inf. Sebastian Deutsch and Dipl. Des. Stefan Landrock developed the basic idea behind Salon when they were given the chance to take over university courses at HFG in Offenbach. Together with their students they built a working prototype of their idea so they could use it for their courses and especially for their presentations. When other universities heard about Salon they were asked if they could host a system for their students too. But Salon was not built to be deployable for other universities and so they had the idea to completely rewrite and to extend the features of Salon so that it could easily be set up for other universities.

3. Implementation

3.1 Technology

The backend of Salon is implemented in Ruby on Rails (short Rails), a web framework written in Ruby¹ and modeled after the MVC software pattern² that allows to quickly

¹ <http://rubyonrails.org/>

² <http://betterexplained.com/articles/intermediate-rails-understanding-models-views-and-controllers/>

create solid web applications without having to care about low-level problems like session-handling or database access. The underlying database is MongoDB³, a document-oriented database system that was chosen because of its flexibility (document-oriented databases are schema free⁴) and its very good integration into Rails.

Salon does not make use of the frontend tools of Rails because the frontend is designed to work as a Single Page Web App (SPWA) and therefore all Rails frontend tools have been replaced with tools that are written in JavaScript so that they could get executed in the browser (see SPWA#intro).

The communication between the frontend and the backend is realized with a REST⁵ interface and all data is being sent in the JSON⁶ format, a format that is very easy to use in both JavaScript (frontend) and Ruby (backend).

3.2 Data Model

[ADD schema image here]

The underlying data structure of Salon is rather simple. There are users that are used for authentication and have basic properties like a username and a password. Pages are associated to users in a one-to-many relationship, which means that users can have as many pages as they want and each page belongs to only one user. Pages have properties like a title, a cover image and a publish state. Each page has a list of assets that are also associated in a one-to-many relationship so that each asset can be associated to one page. Asset is the parent class for image and it stores properties like a title, a link-to location and a position on the canvas.

The reason for deriving image from asset is to allow other assets like for example texts in the future (see salon#quovadis) and to provide all derivations with the needed properties to have a valid asset. The image then only needs to save special properties like the image files and its display sizes.

All assets have a list of tags that are associated in a many-to-many relationship which means that a tag can belong to many assets and assets are able to reference many tags. Tags are used in the search (see salon#pages#search).

³ <http://www.mongodb.org/>

⁴ http://en.wikipedia.org/wiki/Document-oriented_database

⁵ http://de.wikipedia.org/wiki/Representational_State_Transfer

⁶ http://en.wikipedia.org/wiki/JavaScript_Object_Notation

3.3 Navigation

[ADD screenshot of navi here (initial and completely expanded version)]

The navigation in Salon is designed in a breadcrumb-like style⁷. When first visiting the website, the user only sees a caption saying "Salon" which should tell him that he is currently in the most top level of the website. The second element in the navigation is the search bar. Then when the user goes to the overview of a user, the caption "Salon" is replaced by the name of the user that this page belongs to. Removing the "Salon"-caption should emphasize that Salon is about the users and about the work they want to present and that it is not about the platform itself. Normally the first element in a breadcrumb navigation brings the user back to the index page but this is not a scenario that applies to Salon because the index page, intentionally, doesn't offer more features than any other page in Salon but the list of recently created and edited pages. Searching for images is possible from any page through the search field that also resides in the navigation (more on search in [x.y]). If a user wants to go back to the index page he can do this via the menu in the top right corner (more on the Menu in [x.y]) at any time. Another element, the name of the current page, is added to the navigation when the user navigates to a page of a user, and another one, the position of the current image in this set of images, is added when the user navigates to a specific image of a page. The breadcrumb navigation helps the user to keep track of certain information like the owner of the current page and the page an image belongs to. In that way these relations don't need to be displayed on every image or every page, which leads to a cleaner and lighter interface. Also the navigation helps a user to quickly jump back to a user's overview without having to manually navigate there with the back button.

3.4 Menu

The menu is consistently placed in the top right corner of each page and is by default not expanded so that it doesn't unnecessarily take away screen space. To expand the menu the user simply needs to hover over it with the mouse or tap it on the screen (on touch-based devices).

There are two states for the menu: a) The user is logged in; b) The user is not logged in.

[add screenshot of menu when user is not logged in]

When the user is not logged in the menu will have the caption "Sign in/up" which stands for the two most important options that are displayed in the expanded menu. The first point in the menu (see Fig. XX) will lead the user to the sign-in form and the second

⁷ [http://en.wikipedia.org/wiki/Breadcrumb_\(navigation\)](http://en.wikipedia.org/wiki/Breadcrumb_(navigation))

point to the sign-up form. The third point will lead to the about page, that explains the concept of Salon, and the fourth point will lead the user back to the index page.

[add screenshot of menu when user is logged in]

The caption for the menu when a user is logged is its username. The first entry in the expanded menu now leads the user to his own overview page, which allows the user to quickly jump there from any other page. The second entry will lead the user to the page index (more on that here x.x) and the third one will lead the user to the account page where he can edit details of his account.

The fourth entry is highlighted and it is an interactive entry because when this entry gets clicked it will reveal a simple form that lets the user create a new page right from the menu [show screenshot with the states of the form]. After submitting the form the user will get forwarded to the newly created page. By placing the form inside the menu there is no need to create an own page just for adding a new page and the user is able to create pages no matter on what page he is currently on. He only needs to be logged in. The fifth entry triggers a log-out and a redirect back to the index where also the last entry is leading.

3.5 Pages

3.5.1 Index

The index page has, as well as other elements in Salon, two states that depend on the login state of the user. If the user is not logged in the index page displays a text that invites the user to register an account at Salon and a link to the about page so that new users quickly get an idea about what Salon is and how they can use it.

If the user is logged in, the text on the index page welcomes the user and a list of recently created and edited pages is shown at the bottom of the page.

The index page does not have much functionality since the discovery of pages and images is realized with the search field that is located in the navigation (more on search here [x.x]).

3.5.2 Registration and Account

In Salon users have to sign-up with a username, an e-mail address and a password. To complete a registration the user is sent an e-mail with a confirmation link. This step is needed to confirm that a user registered with a valid address. After clicking the link in the mail the system redirects the user to an empty page called "untitled page" and a message is shown that the account now has successfully been confirmed. Since this is

the first time the user uses Salon, a message on the page tells him that he now is able to add images to the page by dragging them onto the window. Also a link to the about page is shown so that new users can get a quick overview on the features of Salon.

The edit account details, like the password or the e-mail address, users can use the "Account" page (link in the menu). This page also allows the user to delete the account.

3.5.3 User Overview

[add screenshot of a simple user overview]

In the user overview all, published and not hidden, pages of a user are displayed on a canvas. Since pages can have a cover image, on this page only the cover images are shown. If a page does not have a cover image, a default picture will get displayed instead. The user is able to arrange all images just by dragging them around (more on DnD here [x.x]). The positions are being saved to the server so that this page will look the same for all visitors and just as the user wants it to look like. Visitors themselves can also drag the images around and create a new layout but the position will not get saved to the server since only the owner has the right to decide how his pages look like.

When dragging an image the image will get populated to the top of all other images so that users can easily create nice effects with occluding images. All these changes will all automatically get saved to the server without the need for the user to initiate the save-process.

To highlight the importance of the images and especially their arrangement there are no further information displayed on top of each image. This is also done so that text elements don't clutter or disturb arrangements that contain a lot of images.

[add screenshot of two hover states (logged-in, not logged-in)]

Further information for a page is displayed on top of the images when a user hovers (or taps) over one of the images. The name of the page and the number of assets that are contained in this page will then fade in and the image gets a half-lucent overlay to highlight which image currently is being hovered. There is the need for the half-lucent overlay because when many pages are placed in the same page it is hard to find out which of the images just has been hovered. Clicking one of the images or its captions will navigate the user to the overview of the page.

When the current user is logged in there will also be additional controls displayed on top of each hovered image. First there is the control to set the size of the image that lets the user choose between four different size options. Then there is a link to the edit page

of the current page that allows a user to quickly edit the page and there is a link to delete the current page. All delete operations in Salon trigger a prompt before actually deleting an element to prevent accidental deletions.

Furthermore the user is not only able to drag each image around but also the whole page which allows to choose a special "starting" point of the canvas that the visitor sees when he first comes to the page. To drag the whole page the user simply needs to drag the background and all other images will get moved accordingly.

3.5.4 Page Overview

At first sight the page overview looks similar to the user overview. The images can freely get dragged around and the title of each image is displayed when the image is hovered. Logged in users also have the ability to directly edit or delete images with the additional captions here. As in the user overview the user is navigated to the image page when he clicks the image or one of the captions on the image.

[add screenshot of images with special icons]

Besides that, there are subtle changes to some of the images. They have special icons that should indicate that they don't link to the image page but to an external page (see Fig.YY e.g. <http://google.com>) or to an another page of this user (e.g. test). (More on cross-references of images in Image#edit[])

Also there is another caption right underneath the navigation that allows the user to quickly jump to edit form of this page.

[add screenshot of upload progress]

Another additional feature is the ability to directly upload pictures to the page by simply dragging picture files from the file system onto the page. A progress dialog will open up that shows the user how many files are left to get uploaded and the images will after the upload directly get added to the page so that the user can work with them on the page right away.

3.5.5 Page Edit Form

[add screenshot of edit form]

On this page the user is able to edit several aspects of a page like its title or its description. Changing the title of a page also leads to the creation of a new URL-slug⁸ for this page, so that the url and the title of a page always correspond. Underneath the normal form there is a listing of all assets that are associated to this page. When

⁸ [http://en.wikipedia.org/wiki/Slug_\(web_publishing\)](http://en.wikipedia.org/wiki/Slug_(web_publishing))

hovering one of these images, new controls to edit and delete the image fade in. Also another option fades in that lets the user set this image as the cover image. When this one gets clicked the image in the normal form automatically changes.

The user is furthermore able to add new images on this page directly by dragging them somewhere onto the page or by opening up the file dialog with the "Add asset" button.

The uploaded images will then automatically appear in the asset list.

Other than on the overview pages the user here has to manually save changes with the buttons that are placed directly under the navigation. There also is a button to cancel the edit form that will remove all changes the user has made and will redirect the user to the page. The third button deletes the page.

3.5.6 Image Overview

On this page the current image is shown in the original size as the user uploaded it. The image is centered horizontally and vertically so that the images' center lies on top of the pages' center. Like on the other pages, the image here can also get dragged around which is handy for images that are bigger than the browser screen so users can see the rest of each image by dragging it around. The position of an image is not saved to the server because the main focus on this page should not lie on a specific arrangement but on the image itself.

By pressing the right- or left key, the user can navigate through the rest of the images of the current page to quickly get an overview over all images.

3.5.7 Image Edit Form

Like in the page edit form, the image edit form first shows the image to make it clear to the user which one of the images he is currently editing. There are basic input fields to change the title (which will also change the URL), the description and there are fields to add Copyright and source information to the image that are important if the user adds an image that has been taken from another website.

Tags can simply get added to the image by typing a tag into the tag field and pressing the return-, the space- or the comma-key. To delete a tag the user simply needs to press the "X"-symbol next to each tag.

A special feature of images is that they can link to three different locations and the user is able to choose a different link for each image. The three radio buttons underneath the label "Links to" specify which location an image points to.

[add mini screenshot for each state]

"Detail" means that the image links to the detail page of an image where the image is shown in original size which is the default link mode of an image (see `image#show` for more info on that).

"Page" means that the image points to another page of the current user. A drop-down menu is shown where the user can choose the page. This allows the user to create connections between Pages and gives users another way to experiment with the website and to express their ideas. Users could create linked-lists of pages that are in some way connected to each other or they could link pages to show a development of a task where each page shows one state.

"URL" means that the image points to a HTTP URL which could be an external URL like an entry in the Wikipedia or it could be another URL from within the Salon website. A scenario could be that users form a group and therefore they create another user. This user then has a page called "Team" where there's an image for each user that links to the user's overview.

Like in the page edit form, all changes that are made in this form need to get confirmed ("Save"-button) and can get discarded ("Cancel"-button).

3.5.8 Page Index

The page index is a list of all the pages the current user owns and it is the page a user is forwarded to after the login.

Each entry in the list stands for one of the user's pages and gives the user an overview on some facts about this page like the current cover image and the number of images stored in it, as well as the name of each page. Furthermore, the page offers the user the ability to quickly delete pages and the ability to directly go to the edit form for each page.

In comparison to the user overview where the user also is able to delete and edit pages, this page also shows all pages that have a status of "hidden" or "not published".

Pages basically can have three states:

- 1) "published": This page is shown in the user's page overview.
- 2) "hidden": This page is invisible and not accessible for other users.
- 3) "not published": This page will not be shown in the user's overview but it is accessible for other users that know the URL.

The "not published"-state is useful when a user is currently working on a page but he wants to show the page to others to get a feedback but he does not want to have this

page appear in his public overview. To change the publish state, a user simply has to choose the new state from the drop-down list and the page automatically gets updated.

3.5.9 Search

[add screenshot of drop down menu]

The search field is located inside the navigation to allow the user to quickly search for images on any page. When typing into the field, the system automatically starts a search for matching tags and provides the user with drop-down. By hitting the "enter" key or by choosing one of the items from the list the user gets redirected to the search result page. Here all images that match the search term will get displayed and a click on them will navigate the user to the corresponding page.

A scenario for the search could be that a prof wants to find all images and the associated pages that this students have put online for an exercise. Students could tag their images with a specific tag so that the prof can find them by searching for it. Also a tag search can be used by students to get inspiration by searching for images on a specific topic.

3.6 Drag and Drop

As described before, the Drag&Drop-Feature is one of the most important distinguishing features of Salon. Therefore there was the need of a good Drag&Drop-Implementation in JavaScript. Most major JavaScript libraries offer Drag&Drop-plugins today and in the beginning of development the most prominent libraries have been tried out (namely jQueryUI, mootools and script.acoul.us). They all worked great and were very feature-rich including UI-Widgets and many abstractions like automatically sortable tables, but they all lacked support for mobile browsers which is an essential feature-requirement for Salon because it should be usable on all iOS devices. Also when using one of the libraries mentioned above one had to include the whole library into the project although only the Drag&Drop functionality was needed. This would add an enormous extra load time especially for users on mobile devices.

Because of these facts I decided to write an own Drag&Drop implementation that supported webkit-mobile browsers as well as desktop browsers. There are basically two ways implement a Drag&Drop System with the given DOM-Events in JavaScript.

Global

This implementation is called to global method because the drag-handler, a component that receives all events and maps them to drag events, stores only one DOM element at a time and associates all "global" DOM events with this element. The drag-handler starts

when the mousedown-event (touchstart on webkit-mobile) is fired on an element with the css class "draggable". This element is then saved as the global drag-target together with its current position. All mousemove-/touchmove events that are then fired on the document initiate a movement-delta calculation and a custom drag-event that is fired on the current drag-target.

Drag-events will get fired until a mouseup or touchend has been fired, which means that the user has stopped dragging an element. This invokes a dragend-event being fired on the drag-target.

Although this method works perfectly on desktop browsers and also on mobile browsers it has some downsides when it comes to touch-device users. When letting iPad users drag elements around a test page they were all confused that they could only drag one element at a time. Also the drag-handler didn't work well when multiple touchstart-events were fired. The fact that element-movement is detected by move events that are fired on the document only allows to track one finger at the same time. Also the iPad users were not only confused but they also thought that the app was not working properly.

Local

To allow multiple elements to get dragged at the same time there was the need to not only associate DOM-events to one single element. Each draggable element now needed its own drag-handler and the global mousemove-events could not be used anymore. Instead of the global events in this system the local move-events are taken to fire drag events. This means that the drag-handler detects drag events from mousemove/touchmove-events that have been fired on this element. The movement delta is calculated not with one global last-position but with a last-position object that is stored for each drag-target.

A problem with this technique is that it does not work well on desktop browsers. When moving the mouse very fast the drag-target lost track of the mousemove-event and the element would stop moving although the mouse was still in movement. Somehow this problem did not appear on touch-devices so that this technique could still be used on touch devices.

The final solution uses both techniques and switches to the global system on desktop browsers and to the local system on touch-devices.

By design the system itself does not alter the positions of the images itself. To make the system as decoupled as possible this functionality has been delegated to event receivers that then can decide on their own in what way they want to move the elements on the screen (e.g. top/left CSS-attributes or negative margins). This makes it possible to create elements that can only get dragged on one axis (horizontal / vertical) or only in a certain range on the screen.

Drop

Although this functionality is not (yet) used in Salon, the Drag&Drop-system also supports the events "drop" and "drag-over". To let an element receive these events it only needs to have the CSS class "draggable" assigned. When a drag event is fired, the system looks for elements that can receive a "drag-over" event by matching the current position with the positions of all draggable elements. This event is useful to give users a feedback that they can drop elements on this element e.g. by increasing its size or by changing its color. If the user drops an element the underlying "draggable"-element will receive the drop event that includes the current drag-target.

3.7 Visual Design

Disclaimer: The visual design has been developed together with Dipl.-Des. Stefan Landrock, Mathias Baer and Dipl.-Inf. Sebastian Deutsch and mainly adopts the style of the Salon prototype.

The main requirement for the visual design was to use as less screen space as possible to not distract the user from the custom design of the pages. Therefore only the navigation, including the search field, and the menu caption are permanently visible.

For the same reason the design itself is very minimalistic and does not use images but is designed with CSS only. To highlight which elements of the page belong to Salon (because users could loose track of that on pages with many images), all elements have been decorated with a drop shadow and are always on top of all other elements of a page. In this way they also remain accessible at any time.

To make it clear to users which elements on the page can only be seen by users that have the right to edit parts of the page, these elements are highlighted in a red-orange color. This also makes it clear to the owner with which of the visual elements he is able to make changes to his pages. The highlight-color is applied to elements like the "edit"-buttons or the "image-size"-widget.

Page owners would normally always see these elements and they would never be able to see their pages as they appear to normal users. For that reason the "Hide admin"-button has been introduced. When clicked it will hide all admin elements so that page owners can get an impression of what their page looks like to normal users.

4 Quo vadis Salon?

The development of Salon should not stop after this thesis and there are various additional features planned for the future:

Remix-me

The "Remix-me" feature would allow a user to clone an existing page from another user to then edit it as if it was one of his pages. These remixes would then get listed on the original page and a caption would get added to pages that are remixes so that the original authors would always be mentioned. This feature should be an optional feature for pages and should need to get activated in the page edit form for each page to preserve copyrights of the original content.

A scenario could be that Profs create pages to specific topics and then make an exercise in class that students should remix the current page and add their own ideas to the page. This raises the question whether it should be allowed to the remixer to delete images from the original author.


Overall this feature could boost interaction between users and could be an element that is fun to use.

More Assets

Currently there is only one asset type that can get added to pages: the image. But the backend design allows to easily add other sorts of assets for example texts or sounds. Especially the combination of text assets and image assets could lead to a lot of interesting pages. Texts could link to other pages or they could serve as a description for images.

Animations

To improve the user experience, more animations could get added to Salon. For example when navigation through the images of a page with the keys, the images could slide-in from the side instead of just suddenly appearing after a key has been pressed.



Or a more complex animation: When selecting a page on a user overview, all other images could fade out. After all other images have been faded out and the new data for the page overview has been fetched, the new images for the page are placed behind the current image. They would then simultaneously fade in and get smoothly moved to their positions while the cover image fades out.

An advantage that comes with this animation is that users are able to see where all the images are placed, even those that would not be visible from the beginning. But this animation could distract the user's attention from what is really important on each page: the images and their arrangement.

Zoom

Pages with a lot of images tend to look cluttered and then images are often dragged out of the visible area of the canvas. Sometimes this leads to several images get hidden by accident so that users will not notice that there are more images on the page. A zoom feature that lets users zoom in and out of overview pages would help to give users an overview of pages with a lot of content and could also be another element for users to experiment with (e.g. by hiding images that can only be seen when a user completely zooms out of the page).

Grid Generator

Sometimes it is not necessary to align images on a page in very creative way when a user only wants to upload the images and show them to someone else. Currently to align them properly in a grid a user has to manually drag the images around. The more images there are, the more time is needed to create a nice grid and very often images are not aligned 100% correct because it is hard to align everything manually. To allow the user to simply create grids that are perfectly aligned, a grid generator component could get added to Salon that is visible on overview pages. A user would only have to specify the amount of columns and the padding and the generator would then align the images automatically.

Search for page titles

The search currently only works for tags but this is a functionality that is likely to get extended in the future to also support the search in page titles and image titles. This makes it possible to find pages directly and not just by searching for tags that images in a specific page may have been tagged with. Students could then name their pages according to the exercises of Profs so that Profs can find these pages easier. This also leads to a needed restructure of the search result page that then should also show resulting pages.

5 Evaluation

In conclusion it can be said that the implementation of Salon went very well and that all features that were available in the prototype are also available in this implementation. This implementation even contains more features than the prototype.

Each page and each image now has a dedicated form to edit its properties, images can have several link targets and the whole website has been implemented as a SPWA which brings a lot advantages in User Experience. Also an own Drag&Drop-system has been written to better match the Salon requirements (more lightweight, (multi) touch-support).

The clean and well-structured code base makes it easy to make changes to the system when universities or customers want to change a behavior or if they want to remove a certain feature. Since Salon is a SPWA with a Rails backend it is also very easy to deploy instances of it to the Internet or a university intranet.

The overall User Experience and the unique Drag&Drop-portfolio concept make Salon a website that is fun to use both for users that create content and for users that browse the page.

II Single Page Web Apps

1. Introduction

1.1 Motivation for Salon

In the beginning Salon was a standard Ruby on Rails (Rails) application. All views were rendered on the server and a lot of JavaScript code was needed to make the UI as flexible as it is now. The JavaScript code was structured with the help of Backbone.js⁹, a

⁹ <http://documentcloud.github.com/backbone/>

JavaScript library that provides frontend developers with Models, Views and Controllers and makes frontend code event-driven. Quickly that lead to duplicated code that needed to get implemented in the backend language and in the frontend language. An example: In order to dynamically create images in the page overview, a JavaScript template was used that looked the same as the ruby template and needed to get maintained in both places. Also parts of the model-logic have been rewritten to enable an easier communication with the backend. More and more of the application-logic moved to the client-side and the code became disordered and it was unclear which parts of the application-logic were implemented in the frontend and which parts in the backend. To restructure and to cleanup the code-base Salon has then been rewritten as a Single Page Web App (SPWA). Besides the code cleanup this also had some positive affects on Salon, which are described in SPWA#Benefits.

1.2 Comparison

Single Page Web Apps gained a lot of attention with JavaScript becoming more and more important in the web development tool chain. One of the main reasons for this development is AJAX¹⁰, a technology to asynchronously load content from the server in JavaScript. Most modern web pages use this technology to make their pages more interactive.

Single Page Web Apps take this approach to a next level by shifting a lot of traditional backend functions to the frontend.

In the following the main differences between the traditional (MVC-based) Web App system and the Single Page Web App system will be pointed out by looking in detail at the typical request flow in both system.

¹⁰ Asynchronous JavaScript and XML

1.2.1 Traditional

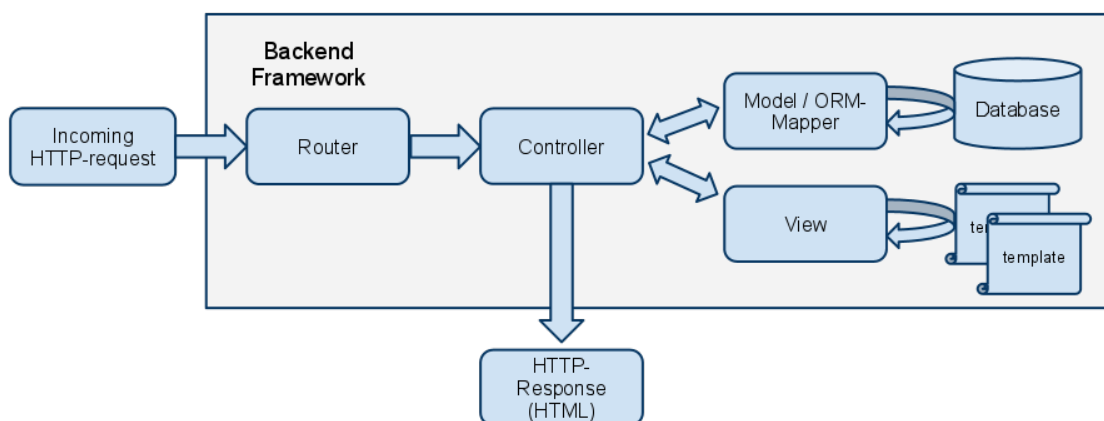


Fig. X: Traditional request flow with a MVC backend

When a new page is requested, the user's browser sends a HTTP request to the server. This request is then matched to the corresponding controller by a router. The controller then triggers the model-layer to retrieve the necessary data for the request from the underlying database. When the data has been successfully fetched, the controller triggers the view-layer to render the data into the requested view. Views may be constructed from several sub-views that are rendered into one layout file. The rendered layout is then transferred back to the browser.

The browser completely removes the old DOM and re-renders a new DOM from the layout that just has been transferred. This leads to the typical blank screen in the browser when a new page is being requested.

1.2.2 Single Page

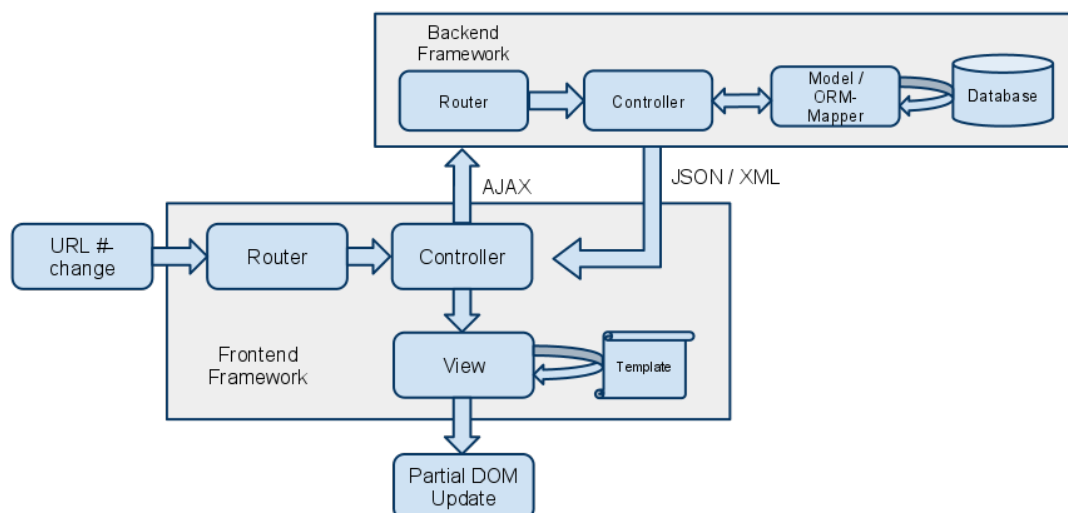


Fig. X: Request flow in a Single Page Web App with a MVC backend and a MVC frontend

The Single Page request-flow is initialized by a "#-change" of the current URL, rather than by a normal URL change (see chapter "[2.1 URLs](#)"). This change triggers a client-side Router to call a certain Controller function that initializes an AJAX request instead of a synchronous HTTP request.

Then the Single Page request-flow is the same as the traditional request-flow but the backend system won't render a HTML view from the data. Instead of letting the server render a complete new layout and transfer it to the client, the data gets serialized into a transport format (e.g. JSON or XML). The frontend framework takes care of rendering a view from the data. Now instead of replacing the whole DOM, the generated view replaces only a part of the page (e.g. the content panel). This prevents the screen from turning blank when the new content is being displayed.

Another difference in the request flow is that in SPWAs the initial request (e.g. "www.mydomain.tld") delivers a complete web app, with all the application-logic in it, and not just a snapshot of it. When the initial page has been loaded, a fronted-router takes care of rendering the correct JavaScript views.

2 Differences

The new structure in the request-flow that has been pointed out in the previous chapter leads to many differences in the development process. These differences are analyzed in this chapter.

2.1 URLs

Since browsers automatically handle hyperlinks with a GET-request, the URLs in SPWAs look different to normal URLs. They make use of the "#"-symbol that originally was used as an anchor to an element in a HTML page. This is needed on pages like Wikipedia where you have long text articles on one page and you want to point the user to a specific paragraph. The browser's viewport automatically jumps to the element with the given ID if there is one.

To prevent the GET-request, a client-side router listens to changes in the URL after the "#"-symbol. [There is no real way to "listen" to changes in the meaning of event-driven programming. This is mostly implemented with polling.] When it changed the router typically triggers a handler for the new URL-partial. So when the URL before was "mydomain.tld/#index" and it has been changed to "mydomain.tld/#about", the URL-router will trigger a handler that has registered for the "about" partial. This also makes

all URLs bookmarkable since the router will render the corresponding views to each URL-partial no matter what page has been displayed before (even on the initial request).

2.2 Rendering

A common technique to render views in the backend is to use an abstraction layer called "templating-engine". These engines allow writing the views in a mostly HTML-like syntax to improve readability and maintainability over manual string-concatenations in the backend language. Also the syntax makes it easy for designers to create and edit templates on their own, rather than having a backend developer who always needs to implement all changes.

A templating-engine pre-compiles views into functions or string-concatenations, so that the backend can execute them faster and doesn't need to interpret them at runtime. The pre-compiled functions are mostly automatically generated and developers generally don't need to work in these files.

In Single Page Web Apps the templating-engine resides on the client-side (see Fig. X). To reduce bandwidth, the backend only sends raw data to the clients (in a transportation format like JSON) that can very rapidly get serialized.

The template-engine in the frontend then renders HTML snippets from the data and inserts them into the DOM.

There are several JavaScript implementations of the most used backend templating-engines and they can all compete with their backend implementations when it comes to speed and flexibility.

2.3 Authentication

State-of-the-art authentication systems are designed to get as easily integrated into a website as possible. Therefore they automatically create views for all authentication actions (sign-up, log-in etc.) that then can get integrated in layout files.

But since in SPWAs the backend views cannot be used, the authentication systems have to get adjusted to support the authentication over AJAX requests. The complexity of these adjustments depend on the modularity and open-ness of the authentication-systems.

If it is not possible to make changes to the authentication system to support AJAX authentication, for example because it's a closed-source project, the complete authentication part could get moved back to the server-side. This means that all forms

that are used for authentication could get served as normal websites instead of client-side views. In this way, all authentication systems in the market could get used.

But a problem of this solution is that there has to be an implementation of the website's layout on the server-side and on the client-side so that all authentication pages look like the rest of the website. This means that layout files have to get maintained in both systems.

2.4 Internationalization

By moving all views to the frontend, also the internationalization (i18n) logic has to be moved to the frontend. I18n systems in modern web application are very well integrated in the view layer because that's where their functionality is primarily needed. But since in SPWAs all views are rendered in the frontend these systems can no longer be used.

There are various i18n systems implemented in JavaScript and they have already proven to work good but they all have a very different API to normal i18n systems so that developers would need to first learn a new API before they could use it. Because of this Salon uses an own i18n system that adopts the API of the Rails i18n system and is therefore easy to use for developers that have already worked with Rails or other backend frameworks.

One disadvantage of client-side i18n systems is that all languages have to get loaded in the initial page request to have a smooth and quick language-change. When it is not important to have a new language to load immediately, the translations files could also be loaded when they are needed. Mostly the language is not so often changed by a user, which means that loading the language files only when they are needed seems like a valid method.

2.5 Notifications

Many backend frameworks give developers a simple way to display so called flash messages, which are short messages that are displayed once a request has been finished and the page has been rendered, like "Successfully deleted this item".

The purpose of these messages is to give the user a feedback to remind him what he just did on the previous page and whether this action has been applied successfully or whether an error occurred during the processing. A typical backend framework provides these messages to the view layer where the messages normally are being rendered into a DOM element.

This technique doesn't work for SPWAs since the backend doesn't render parts of the view so there is the need for a new notification system in SPWAs. A simple system has been implemented for Salon to provide Salon with an easy to use notification system that is also easy to use in other SPWAs projects. The system adds the flash messages to the JSON response on the server-side and a client-side notification component, which listens for all incoming AJAX responses, parses the message and displays it accordingly (distinguishes between success and error messages).

Notifications and status indicators are very important for SPWAs and other AJAX-heavy websites because there is no reload of the page that tells the user that something is happening on the page. AJAX requests may take a long time so that users should also get an immediate visual feedback to show the user that the website is working on his input. And since in most cases only parts of the website change there should be notifications that tell the user exactly what just has happened because the user may not notice minimal changes in the page layout.

2.6 Forms

One of the things that enormously speed up the development when working with a backend framework are form-helpers. They allow developers to rapidly create forms for CRUD [explanation, CREATE;RETRIEVE;UPDATE;DELETE] operations on models without having to write much code. Again, these helpers don't work in SPWAs since they generate forms in the backend and create HTML that is used in backend templating-engines.

Writing all forms without such helpers would lead to a lot of duplicated code that is needed for the AJAX requests and the error handling and it is also very time consuming having to write forms from scratch each time. To avoid the unnecessary time consumption and the duplicate code a client-side form-helper has been developed during the development of Salon. The component works well together with the model base classes that have been provided by the already mentioned Backbone.js (see XX) and automatically takes care of the server communication and validations. The form-helper is additionally capable of pre-filling forms with a model's attributes, it provides hooks to allow developers to override the default behavior (e.g. AJAX calls) and it automatically displays error messages when a user has entered wrong values.

3 Benefits

3.1 Speed / Efficiency

The most important benefit of SPWAs is that they'll speed up your website performance. They will even make the client-server communication more efficient. These benefits are caused by sending less data to the clients (no HTML, just transport format) and by the less amount of time that is needed by the server to render complex views.

Efficiency is very important for example when a lot of clients connect to a website via slower networks or when a website server will have to handle a lot requests per second. The faster a website reacts on user input or the faster it loads, the better is its user experience.

There are a lot of studies that investigated the impact of a website's speed on its user experience and they all support the thesis mentioned above. For example, in 2009, Forrester Consulting conducted a study to investigate the behavior of online shoppers¹¹. They found out that a page should not take longer than 2 seconds to load or otherwise the user becomes unsatisfied and eventually will stop using the online shop or the user may even switch to another competitor. 52% of the interviewees mentioned in the poll that page speed is one of the most important features for a good online shop.

When Google intentionally slowed down their search results in one of their public experiments¹², they observed a decline of the total number of searches by 0.2% to 0.6%. They found out that the more delay they added to the results, a user would make less searches. When considering the shortness of the delays that Google added to the searches (first 100ms, later up to 400ms), this experiment very well shows how important each millisecond delay can be for the overall user experience on a website. The simple and efficient design of client-server communication in SPWAs makes them very fast so that the time users have to wait is reduced to a minimum.

3.2 Same Language

Server-side JavaScript has become very popular recently with the development of node.js, an event-driven server that allows developers to write all their backend code in JavaScript. Its event-based programming paradigm, I/O operations won't block the server until they're finished, instead an event is fired when data is available, allows the server to handle way more concurrent request than other (blocking) server technologies. This makes this technology very interesting for high-traffic websites.

¹¹ http://www.akamai.com/html/about/press/releases/2009/press_091409.html

¹² <http://googleresearch.blogspot.com/2009/06/speed-matters.html>

Dealing with the same language on both end-points means that developers can share code, for example model-logic, to reduce code duplication and unwanted double-maintenance. This can lead to a faster and more efficient development and developers don't need to change the way they program since it is the same syntax in the frontend and in the backend.

3.3 User Experience

Talking about speed in the context of user-experience means more than just performance of client-server communication¹³. When a website does heavy calculations for the user SPWAs won't perform better in calculating the result on the server-side than normal websites. But one weakness of normal websites is that there is no feedback telling the user that a certain operation may take longer. The screen will stay blank or will not change for a long time.

In SPWAs, loading indicators like labels (e.g. "Loading...") or spinning animations are used to give the user an immediate feedback on an action that may take longer. This won't speed up the calculation but it shows the user that the system has registered the input and that the user has to wait. This also prevents users from clicking the same link again, which may even lead to longer response times.

Transitions

Another benefit SPWAs have over normal web pages is that it is possible to have transitions / animations between page changes.

These effects are more and more often used in modern web pages to make them feel more dynamic and to make the user have more fun using the page. But a reload on a normal website will break the dynamic impression because the page will simply turn blank on page change until the new page is loaded.

A good example for this behavior is Apple's website. When visiting the index page only a main focused element, like for example a new iPad, is shown. All other elements like the navigation fade in after a certain delay. But when going to another page, the page turns blank and the new page is being loaded which completely breaks the smooth impression of the index page.

To make the User Experience on a website consistent, page transitions like they are implemented in the Salon overview pages could get added so that the dynamic and interactive impression will not get lost. On Salon overview pages all images fade-out and

¹³ <http://code.google.com/intl/de/speed/articles/usability-latency.html>

fade-in when navigating through the different pages. This makes the navigation feel a lot smoother and it also hides loading times (caused both from the AJAX request and the image loading time).

Users normally don't actively see the transitions. It is, when it's done right, only a subtle effect. But when transitions are removed from a page, users see that something has changed and that the page does not feel as good as before. Showing users first the transitioning implementation of the overview pages and then an implementation without transitions has proven this. Users thought that the site needed longer to load but they could not tell that the transitions had been removed.

Transitions should not be used too heavily because otherwise the focus in the important parts of a website will get lost and users tend to get annoyed by too long animations.

Sound

Another neat feature that SPWAs offer is that they allow websites to have music playing permanently in the background without stopping when the page changes.

Currently most websites that let users play music suffer from the page-reload problem and don't allow the user to simultaneously browse the page and listen to the music (e.g. <http://www.last.fm>, <http://www.soundcloud.com>). Users have to keep at least two tabs/windows of these pages open.

Other websites bypass this problem by opening a dedicated new window only for the music player (e.g. <http://www.jamendo.com>, <http://www.play.fm>). Both solutions suffer from the same problem: it is very cumbersome for the user to control the player. The user has to switch the tab / window or even, when the user has too many tabs / windows open, the user has to search for the player.

Stopping the player or altering the volume can take quite a while and this delay leads to a bad user experience. In SPWAs the player can simply get embedded into the page and it will always remain on the same position so that users can easily control it. Of course the player won't stop playing when another page is requested since the reload does not affect the player.

A good example for the use of SPWAs in a music-context is simfy (<http://www.simfy.de>) [add screenshot, maybe with comparison to other sites mentioned above]. The player is fixed at the bottom of the page and it remains there when the content of the page changes. To not get in the way while browsing, the player's size is adjustable.

4 Problems

4.1 Search Engine Optimization

Search Engine Optimization is very important for modern websites to get a good ranking in search results from Google or any other search engine. Search engines build their indexes with so called Web crawlers¹⁴ that process the contents of websites to get an understanding to what topics they are related. Web crawlers automatically follow links on web pages to create relations between websites and to find out the importance of websites by counting the links that link back to a certain page.

They are built to rapidly crawl through many websites, which means that the basic crawlers neither load images nor CSS-, nor JavaScript files to decrease the loading time. This has a negative impact on SPWAs because the content would not be correctly indexed or even not indexed at all because the client-side JavaScript-based URL-router would not get started when a Web crawler is on the website, since they don't execute JavaScript.

Furthermore, if the start page of a SPWA was generated by the JavaScript templating system, the page wouldn't even get added to any search engine index because the crawler would just see a blank HTML page. This would make SPWAs for now unusable for projects that need to have a good ranking in search engines.

But there are several ways to go around this problem.

Google proposed a technique that let's their crawlers index an SPWA¹⁵.

When their crawler finds a URL in the SPWA-typical "#!"-style it will request a special URL on the website's server that should return a HTML snapshot of the requested page. This snapshot should represent the content of the page that should get indexed. As an example, a request to the URL "mydomain.tld/#!/test" would create a Web crawler request to "mydomain.tld/?_escaped_fragment_=test" and the server should respond with the HTML snapshot.

This solution can easily lead to a lot of duplicated code since a router is needed in the backend that has to work exactly the same as the router in the frontend. Additionally, often view code might get duplicated because templates have to get maintained in the frontend and in the backend.

¹⁴ http://en.wikipedia.org/wiki/Web_crawler

¹⁵ <http://code.google.com/intl/de/web/ajaxcrawling/docs/getting-started.html>

One thing to keep in mind with this technique is that currently only the Google Web crawler supports the advanced URL scheme and none of the other competitors such as Bing and Yahoo do.

Another technique works especially well for community pages where there is a difference between the site a user sees when he is logged in and the site he sees when he is not logged in (Web crawlers normally don't log in to communities). In that case all public pages could get served from the backend so that Web crawlers easily can index them. Internal pages would then get served as SPWA since these pages should not get indexed anyway. The extra effort that is needed for this technique is reasonable since only few pages (e.g. index, about, pricing, contact) need a backend view and most of the client-side code doesn't need to get duplicated.

4.2 New Tools needed

When working on a normal web application there are tons of frameworks and tools that help throughout the whole development-, deployment- and maintenance process. These tools have been optimized over the past years and developers have learnt how to become most productive with these tools.

Such integrated tools and frameworks don't exist for SPWAs yet. There are tools that fit one specific part of the process like compiling the Views or giving the app a MVC structure, but developers have to connect these tools manually. In the backend the old tools can still get used but they don't support the client-side development that much. New tools have to get developed so that developers don't have to struggle with their development environment on every new project and can clearly focus on working on the project.

A first step for new tools has been made with brunch¹⁶, a tool-chain that combines all the needed technologies for the client-side development into one command line call. This very much helps to speed up the development on SPWAs but developers still have to develop the back-end with another tool because brunch currently is backend-agnostic and doesn't provide any backend helpers. To ease development even more there is definitely the need for tools that support the development in the frontend as well as in the backend.

¹⁶ <https://github.com/brunch/brunch>

4.3 Exposure of Business Logic / Sensitive Data

Moving the complete business-logic to the client-side means that every user that knows how to display the source of a website can easily see how a website / a business works. Modern browsers even further have integrated tools that allow users to deeply inspect a website's code and to especially monitor AJAX requests. The problem that users can cheat on a website's code is not a problem that only occurs in SPWAs, generally a problem of every website, but since almost all business-logic resides in the user's browser this problem has a much bigger impact on SPWAs. When a website deals with sensitive data (bank accounts, credit card numbers...) developers need to make sure that none of the code on the front-end exposes security holes that could harm the website's users. Generally all privacy relevant operations should be done on the server-side and to secure the client-server communication HTTPS should be used instead of normal HTTP. Another important thing is to double-check login-states and admin-rights on the server and don't let only the client-side handle it. Assume the following scenario: An app has a global user object with a boolean field called "admin". The app displays editable elements according to the value of the admin field in the user object. A user could now simply open up the JavaScript console of his browser and change the value of the admin field to get access to all editable elements.

A way to protect websites from this attack is to add a server-side generated field in the responses that adds information on the rights of the current user. The app then switches the editable elements according to the response rather than according to a front-end object that could have been edited by a user. But still the backend should always double-check the user's rights on incoming requests.

Another method to secure a website's code is to obfuscate it before it gets deployed to the server. In that way the code is not as readable as before and attackers would need to put a lot more effort in understanding the obfuscated code code to harm a website. A nice side effect that usually comes with code-obfuscation in JavaScript is that the code also gets compressed so not only the website gets secured but also the load time gets reduced.

4.4 Accessibility

SPWAs only work when clients have JavaScript enabled and this makes them useless for users that either have disabled JavaScript on purpose, e.g. to increase the browser performance, or users that are browsing at work and that are not allowed to have their browsers execute JavaScript.

According to statistics by Yahoo¹⁷, the percentage of users that have disabled JavaScript is approximately 1%. That number may seem low at first but it means that if a page has one million page-visits a month, ten thousand of them are not able to use the page.

There is the possibility to display a message that demands the user to activate JavaScript with the “noscript-tag” in HTML¹⁸, but users that have actively disabled JavaScript did this for a reason and may simply decide not to use the site. Developers have to be aware of the numbers mentioned above when creating JavaScript-heavy applications.

Another problem with dynamically generated layouts is that it is hard for screen readers to semantically interpret these layouts. Although 75% of screen reader users have JavaScript enabled¹⁹, the screen readers are not able to properly interpret dynamic DOM changes that are used to “switch” pages in SPWAs.

The Web Accessibility Initiative (WAI), an organization that creates recommendations for web developers to make websites more accessible, is aware of the problems described above and created a guideline called “WAI-ARIA”²⁰ that aims to provide solutions to developers and its publication is expected for the middle of 2011. The current working draft (<http://www.w3.org/TR/wai-aria-implementation>) describes a set of additional events that should get fired for example when a DOM element's text-content changes or an element is being hidden. These events are processed by screen readers to create a semantically correct feedback for the user.

5. Conclusion

SPWAs really can help to make a website feel better, to give the user a better experience browsing it. They allow a lot of new interaction concepts and more dynamic sites than we have today. AJAX was a first step to make websites feel more fluid but SPWAs bring the whole concept to a next level by giving the ability to get a completely fluid navigation and transition system. There now is the possibility to create websites that don't look and feel like normal websites and actually are fun to wrk with.

The concept has already been taken over by big companies like Google (Google Mail Chat / Client-side routing) or Facebook (Facebook Chat / Facebook Messages / Content is replaced inline, no new request) and others to make parts of their website more

¹⁷ <http://developer.yahoo.com/blogs/ydn/posts/2010/10/how-many-users-have-javascript-disabled>

¹⁸ <http://www.w3.org/TR/html4/interact/scripts.html#h-18.3.1>

¹⁹ <http://webaim.org/projects/screenreadersurvey2/#javascript>

²⁰ <http://www.w3.org/WAI/intro/aria.php>

dynamic and I think that there will be more and more pages that take over the technique.

In my opinion the rise in the interest for Node.js will also result in more and more companies switching to the SPWA idiom not just because of its positive impact on the User Experience but also because of the ability to share code between client and server. And since there are (currently) no big web frameworks like RoR for Node.js developers it may be easier and faster for developers use tools like brunch to program websites with a simple REST-based Node.js backend.

But one also has to admit that SPWA brings a lot of new tools and technologies with it and so it might in a first run not be faster to develop because the programmers need to learn the new tools. Especially when developers don't have experience in programming in JavaScript the learning process can take a while because of the asynchronous parts that you have to deal with in JavaScript.

Before developing a SPWA one should first think about the negative points above and decide if it's okay to not get ranked in search engines currently (without extra effort) or if it's okay to expose the business logic to the user. Only if you don't think that these negative impacts will harm the success of your website you can start building the app. If you decide to port an already running and established website you should think twice about switching because your page may already be well indexed in the major search engines and all links may lead to a dead end after the rewrite of your page (if you don't have a redirect component) like it happened to gawker.com, a well established "media news and gossip" blog²¹, that switched the whole blog to a SPWA in February of 2011 and all indexed links were broken so that the number of unique visitors dropped by 50%²² and a lot of users wrote bad reviews about the page. So before your switch, make sure you don't break indexed URLs.

²¹ <http://en.wikipedia.org/wiki/Gawker>

²² <http://techcrunch.com/2011/02/17/gawker-redesign/>